Congratulations! You passed!

Grade received 80%
Latest Submission Grade 80%
To pass 80% or higher

Go to next item

1. Consider using this encoder-decoder model for machine translation. 1 / 1 point

This model is a "conditional language model" in the sense that the encoder portion (shown in green) is modeling the probability of the input sentence x.

- True
- False

Expand

Correct

2. In beam search, if you increase the beam width B, which of the following would you expect to be true? 1 / 1 point

- Beam search will run more quickly.
- Beam search will use up less memory.
- Beam search will converge after fewer steps.
- Beam search will generally find better solutions (i.e. do a better job maximizing $P(y \mid x)$).

Expand

Correct
As the beam width increases, beam search runs more slowly, uses up more memory, and converges after more steps, but generally finds better solutions.

3. In machine translation, if we carry out beam search without using sentence normalization, the algorithm will tend to output overly short translations. 0 / 1 point

- True
- False

Expand

Incorrect

4. Suppose you are building a speech recognition system, which uses an RNN model to map from audio clip x to a text transcript y. Your algorithm uses beam search to try to find the value of y that maximizes $P(y \mid x)$. 1 / 1 point

On a dev set example, given an input audio clip, your algorithm outputs the transcript $\hat{y} =$ "I'm building an A Eye system in Silly con Valley.", whereas a human gives a much superior transcript $y^* =$ "I'm building an AI system in Silicon Valley."

According to your model,

$$P(\hat{y} \mid x) = 1.09 * 10^{-7}$$

$$P(y^* \mid x) = 7.21 * 10^{-8}$$

Would you expect increasing the beam width B to help correct this example?

- No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.
- No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.
- Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.
- Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.

Expand

Correct

5. Continuing the example from Q4, suppose you work on your algorithm for a few more weeks, and now find that for the vast majority of examples on which your algorithm makes a mistake, $P(y^* \mid x) > P(\hat{y} \mid x)$. This suggests you should focus your attention on improving the search algorithm. 0 / 1 point

- False.
- True.

Expand

Incorrect

6. Consider the attention model for machine translation. 1 / 1 point

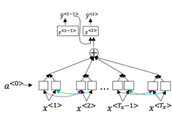

Further, here is the formula for $\alpha^{<t,t'>}$.

$$\alpha^{<t,t'>} = \frac{\exp(e^{<t,t'>})}{\sum_{t'=1}^{T_x} \exp(e^{<t,t'>})}$$

Which of the following statements about $\alpha^{<t,t'>}$ are true? Check all that apply.

- $\sum_{t'} \alpha^{<t,t'>} = 0$
- $\alpha^{<t,t'>}$ is equal to the amount of attention $y^{<t>}$ should pay to $a^{<t'>}$
- We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t'>}$ that are highly relevant to the value the network should output for $y^{<t>}$. (Note the indices in the superscripts.)
- $\sum_{t'} \alpha^{<t,t'>} = 1$

Expand

Correct
$$\alpha^{<t,t'>}$$ = amount of attention $$y^{<t>}$$ should pay to $$a^{<t'>}$$

7. The network learns where to "pay attention" by learning the values $e^{<t,t'>}$, which are computed using a small neural network: 1 / 1 point

We can replace $s^{<t-1>}$ with $s^{<t>}$ as an input to this neural network because $s^{<t>}$ is independent of $\alpha^{<t,t'>}$ and $e^{<t,t'>}$.

- True
- False

Expand

Correct
We can't replace $$s^{<t-1>}$$ with $$s^{<t>}$$ as an input to this neural network. This is because $$s^{<t>}$$ depends on $$\alpha^{<t,t'>}$$ which in turn depends on and $$e^{<t,t'>}$$; so at the time we need to evaluate this network, we haven't computed $$s^{<t>}$$.

8. Compared to the encoder-decoder model shown in Question 1 of this quiz (which does not use an attention mechanism), we expect the attention model to have the greatest advantage when: 1 / 1 point

- The input sequence length T_x is large.
- The input sequence length T_x is small.

Expand

Correct

9. Under the CTC model, identical repeated characters not separated by the "blank" character (_) are collapsed. Under the CTC model, what does the following string collapse to? 1 / 1 point

__c_oo_o_kk___b_ooooo__oo__kkk

- cokbok
- coookkboooooookkk
- cook book
- cookbook

Expand

Correct

10. In trigger word detection, if the target label for $z^{<t>}$ is 1: 1 / 1 point

- Someone has just finished saying the trigger word at time t.
- There is exactly one trigger word.
- The total time that the trigger word detection algorithm has been running is 1.
- Only one word has been stated.

Expand

Correct
Target labels indicate whether or not a trigger word has been said.